

Redis Cluster搭建

Redis集群安装

1. 下载Redis源码

```
tar zxvf redis-X.X.X.tar.gz
```

2. cd redis-X.X.X

```
make
```

3. make install

```
ls /usr/local/bin/ #确认有以下文件
```

```
redis-check-aof redis-cli redis-server
```

```
redis-benchmark redis-check-rdb redis-sentinel
```

```
cp src/redis-trib.rb /usr/local/bin/
```

安装Ruby的依赖

```
yum install ruby rubygems
```

```
gem install redis
```

确认运行redis-trib.rb 不报错

4. 源码目录utils下的脚本建议在今后的工具去阅读一下

5. 计划搭建一个Cluster结构

```
/data/rcluster/
```

```
|—— tpl.sh
```

```
|—— start.sh
```

```
|—— stop.sh
```

```
|—— 7000
```

```
|   |—— 7000.conf
```

```
|—— 7001
```

```
|   |—— 7001.conf
```

```
|—— 7002
```

```
|   |—— 7002.conf
```

```
|—— 7003
```

```
|   |—— 7003.conf
```

```
|—— 7004
```

```
|   |—— 7004.conf
```

```
|—— 7005
```

```
|—— 7005.conf
```

tpl.sh

```
cat <<EOF
bind 192.168.11.101
dir /data/rcluster/$port
port $port
cluster-enabled yes
daemonize yes
cluster-config-file nodes$port.conf
```

```
cluster-node-timeout 5000
appendonly yes
cluster-require-full-coverage no
logfile "./$port.log"
EOF
```

init.sh

```
#!/bin/bash
for ((i=0; i<6; ++i))
do
mkdir 700$i
export port=700$i
sh tpl.sh >>700$i/700$i.conf
done
```

stop.sh

```
#!/bin/bash
for ((i=0; i<6; ++i))
do
echo 700$i
/usr/local/bin/redis-cli -p 700$i shutdown
done
```

6. 初始化整个集群

因为现在有6个节点可以建出来3组一主一从的结构，命令如下：

```
redis-trib.rb create --replicas 1 192.168.11.101:7000 192.168.11.101:7001 192.168.11.101:7002 192.168.11.101:7003
192.168.11.101:7004 192.168.11.101:7005
```

>>> *Creating cluster*

>>> *Performing hash slots allocation on 6 nodes...*

Using 3 masters:

192.168.11.101:7000

192.168.11.101:7001

192.168.11.101:7002

Adding replica 192.168.11.101:7003 to 192.168.11.101:7000

Adding replica 192.168.11.101:7004 to 192.168.11.101:7001

Adding replica 192.168.11.101:7005 to 192.168.11.101:7002

M: 2ed089f797a407d1e20b9f63b39a09b53154df8c 192.168.11.101:7000
slots:0-5460 (5461 slots) master

M: 2128addbc4eaf4566df699a7e98453c0779f6233 192.168.11.101:7001
slots:5461-10922 (5462 slots) master

M: 5d1b8642594618adf310f3086bf50eb23274a6ec 192.168.11.101:7002
slots:10923-16383 (5461 slots) master

S: 6c4a627fee4194905b7c2ed7bdbad137e1ea67ad 192.168.11.101:7003
replicates 2ed089f797a407d1e20b9f63b39a09b53154df8c

S: acc936d980b59369186f3ba3a16d75c01ba7af4d 192.168.11.101:7004
replicates 2128addbc4eaf4566df699a7e98453c0779f6233

S: 5d4900c21d6d0f874951b3b5a5c4ed57d6399bff 192.168.11.101:7005

```

replicates 5d1b8642594618adf310f3086bf50eb23274a6ec
Can I set the above configuration? (type 'yes' to accept): yes
>>> Nodes configuration updated
>>> Assign a different config epoch to each node
>>> Sending CLUSTER MEET messages to join the cluster
Waiting for the cluster to join..
>>> Performing Cluster Check (using node 192.168.11.101:7000)
M: 2ed089f797a407d1e20b9f63b39a09b53154df8c 192.168.11.101:7000
slots:0-5460 (5461 slots) master
1 additional replica(s)
M: 5d1b8642594618adf310f3086bf50eb23274a6ec 192.168.11.101:7002
slots:10923-16383 (5461 slots) master
1 additional replica(s)
M: 2128addbc4eaf4566df699a7e98453c0779f6233 192.168.11.101:7001
slots:5461-10922 (5462 slots) master
1 additional replica(s)
S: 6c4a627fee4194905b7c2ed7bdbad137e1ea67ad 192.168.11.101:7003
slots: (0 slots) slave
replicates 2ed089f797a407d1e20b9f63b39a09b53154df8c
S: acc936d980b59369186f3ba3a16d75c01ba7af4d 192.168.11.101:7004
slots: (0 slots) slave
replicates 2128addbc4eaf4566df699a7e98453c0779f6233
S: 5d4900c21d6d0f874951b3b5a5c4ed57d6399bff 192.168.11.101:7005
slots: (0 slots) slave
replicates 5d1b8642594618adf310f3086bf50eb23274a6ec
[OK] All nodes agree about slots configuration.
>>> Check for open slots...
>>> Check slots coverage...
[OK] All 16384 slots covered.

```

以上输出可以看出来三组主从结构初始化，另外16384个Slot在三组机器上的分布。

到本步骤集群搭建完毕，接下来进行一些测试

集群状态查看：

redis-cli -c -h 192.168.11.101 -p 7000 cluster nodes

```

5d1b8642594618adf310f3086bf50eb23274a6ec 192.168.11.101:7002 master - 0 1499839043346 3 connected 109
23-16383
2128addbc4eaf4566df699a7e98453c0779f6233 192.168.11.101:7001 master - 0 1499839043851 2 connected 546
1-10922
6c4a627fee4194905b7c2ed7bdbad137e1ea67ad 192.168.11.101:7003 slave 2ed089f797a407d1e20b9f63b39a09b531
54df8c 0 1499839043346 4 connected
acc936d980b59369186f3ba3a16d75c01ba7af4d 192.168.11.101:7004 slave 2128addbc4eaf4566df699a7e98453c077
9f6233 0 1499839042336 5 connected
2ed089f797a407d1e20b9f63b39a09b53154df8c 192.168.11.101:7000 myself,master - 0 0 1 connected 0-5460
5d4900c21d6d0f874951b3b5a5c4ed57d6399bff 192.168.11.101:7005 slave 5d1b8642594618adf310f3086bf50eb232
74a6ec 0 1499839044355 6 connected

```

1. 测试

利用redis-cli命令需要带上 -c 参数 (Enable cluster mode (follow -ASK and -MOVED redirections))

```
#redis-cli -h 192.168.11.101 -p 7000
```

```
192.168.11.101:7000> set foo bar
(error) MOVED 12182 192.168.11.101:7002
192.168.11.101:7000> get foo
(error) MOVED 12182 192.168.11.101:7002
192.168.11.101:7000> quit
```

```
#redis-cli -c -h 192.168.11.101 -p 7000
```

```
192.168.11.101:7000> set foo bar
-> Redirected to slot [12182] located at 192.168.11.101:7002
OK
192.168.11.101:7002> get foo
"bar"
```

在这里可以看出来如果没-c参数连接进去，是不支持跳转，直接报错，加上 -c 标识，即可以实现集群间成员的自动跳转。

场景二：

写100个数据观察三个节点上的分布情况

```
for ((i=0; i<100; ++i)); do redis-cli -c -h 192.168.11.101 -p 7000 set zst$i "The Best MySQL edu at zhishutang.com :)"; done
```

```
redis-cli -c -h 192.168.11.101 -p 7002
```

```
192.168.11.101:7002> get zst2
"The Best MySQL edu at zhishutang.com :)"
192.168.11.101:7002> get zst1
-> Redirected to slot [2409] located at 192.168.11.101:7000
"The Best MySQL edu at zhishutang.com :)"
192.168.11.101:7000> get zst3
-> Redirected to slot [10539] located at 192.168.11.101:7001
"The Best MySQL edu at zhishutang.com :)"
192.168.11.101:7001> get zst4
"The Best MySQL edu at zhishutang.com :)"
192.168.11.101:7001> get zst5
-> Redirected to slot [2541] located at 192.168.11.101:7000
"The Best MySQL edu at zhishutang.com :)"
```

1. 集群节点重启

参考stop.sh 如果不行就kill吧。

2. 节点管理

添加一个节：

```
redis-trib.rb add-node ip:port 192.168.11.101:7000
```

添加节点分成添加主节点和从节点，节点最好是空的。第二个参数是集群里成员。

如添加一个主节点到集群：

```
redis-trib.rb add-node 192.168.11.101:7006 192.168.11.101:7000
```

```
> \>>> Adding node 192.168.11.101:7006 to cluster 192.168.11.101:7000
> \>>> Performing Cluster Check (using node 192.168.11.101:7000)
```

```

> M: 2ed089f797a407d1e20b9f63b39a09b53154df8c 192.168.11.101:7000
> slots:0-5460 (5461 slots) master
> 1 additional replica(s)
> M: 5d1b8642594618adf310f3086bf50eb23274a6ec 192.168.11.101:7002
> slots:10923-16383 (5461 slots) master
> 1 additional replica(s)
> M: 2128addbc4eaf4566df699a7e98453c0779f6233 192.168.11.101:7001
> slots:5461-10922 (5462 slots) master
> 1 additional replica(s)
> S: 6c4a627fee4194905b7c2ed7bdbad137e1ea67ad 192.168.11.101:7003
> slots: (0 slots) slave
> replicates 2ed089f797a407d1e20b9f63b39a09b53154df8c
> S: acc936d980b59369186f3ba3a16d75c01ba7af4d 192.168.11.101:7004
> slots: (0 slots) slave
> replicates 2128addbc4eaf4566df699a7e98453c0779f6233
> S: 5d4900c21d6d0f874951b3b5a5c4ed57d6399bff 192.168.11.101:7005
> slots: (0 slots) slave
> replicates 5d1b8642594618adf310f3086bf50eb23274a6ec
> [OK] All nodes agree about slots configuration.
> \>>> Check for open slots...
> \>>> Check slots coverage...
> [OK] All 16384 slots covered.
> \>>> Send CLUSTER MEET to node 192.168.11.101:7006 to make it join the cluster.
> [OK] New node added correctly.

```

添加从节点:

```
redis-trib.rb add-node --slave 192.168.11.101:7006 192.168.11.101:7007
```

说明: 这种情况没有指定是给那个主节点添加从, 整个集群会随机的选出来一个主。新加入的节点上面没有任何数据 (slot) 是一个空节点。

如果需要指定给某个主节点添加从节点, 可以用: master-id指定

```
redis-trib.rb add-node --slave --master-id 3c3a0c74aae0b56170ccb03a76b60cfe7dc1912e
192.168.11.101:7006 192.168.11.101:7007
```

移除节点:

```
redis-trib.rb del-node 192.168.11.101:700 'node-id'
```

如: redis-trib.rb del-node 192.168.11.101:700 '3c3a0c74aae0b56170ccb03a76b60cfe7dc1912e'

```
redis-trib.rb del-node 192.168.11.101:7000 '49bf78b8b0f0cb085fd36cafbcf65cf849a47caa'
```

```

> >>> Removing node 49bf78b8b0f0cb085fd36cafbcf65cf849a47caa from cluster 192.168.11.101:7000
> >>> Sending CLUSTER FORGET messages to the cluster...
> >>> SHUTDOWN the node.
> >>>

```

需要注意: 第一个连接节点是任何一个Redis都可, 后面的node-id是节点的id, 对主节点如果需要移除, 需要把上在贩Slot先迁移走。

利用 cluster nodes 去检验

1. Slot管理

数据重新平衡

新加入节点后, 需要做数据平衡

```
redis-trib.rb reshard 192.168.11.101:7000
```

Step1: 提示数据向那个节点平衡提供相应主节点的node-id

Step2: 指定移走多少Slot

Step3: 指定多那个master上移走, 提供node-id

Step4: 输出done

迁移开始

```
** redis-trib.rb rebalance 192.168.11.101:7000**
```

[illegible]

可以尝试把一组机器里的一个节点关掉试试，或是整个Group里的成员全部关掉。测试一下。

- 若无特别声明，本培内容是知数堂内部教材， 版本归知数堂训训；
- 我们欢迎知识共享，但鄙视不尊重他人的劳动成果的行国，比如直接贴到Blog或是申请原创
- 直接报名参数堂的优势：
 - 由多位资深行业专家亲自授课
 - 课程中可以有效互动，有疑问得到及时的解答
 - 可以和老师以及同学建立行业圈子，更利于职业发展
 - 和一同学习的同学，更专注和系统化的提高
 - 学成后可以获量老师直接推荐到各大互联网公司的机会